

Problem-solving

(For small interview questions)

Terence Parr
MSDS program
University of San Francisco

It's all about the process

- Technical interviews usually consist of small bite-sized questions; that gives us an opportunity to develop a process to solve them
- You might have a successful interview even if you can't solve a specific problem, if you demonstrate a solid problem-solving process
- Multiple interviewers have told me that they are mostly interested in **how** candidates solve problems
- A common tactic, in fact, is to provide an ill-specified or incompletely specified problem
- The interviewer wants to see how you clarify and nail down the actual problem

Overall process

1. Understand
2. Solve
3. Code (notice this is a separate step from #2)
4. Verify/test



1. Understand

- Read the description (3x) then restate the problem, either on paper or out loud (yes, I talk to myself in my office)
- Ask the interviewer if you've understood it correctly
"What I hear you saying is ..."
- Describe and write out a minimal but *representative* example of
 - the intended input data or data structure *and*
 - the expected output *and*
 - ask if the example is correct
- Identify any edge cases you can think of by example, but don't focus on those cases initially

act → *cat*
its → *sit*

Reading problem descriptions

- Details matter, pay careful attention to the interviewer
 - Pretend that they are trying to trick you with the problem description!
 - Are the input data elements strings, ints, floats?
 - If data is numeric, are they always between 0 and 1? Can they be negative?
 - Is the input sorted?
 - Can you see all of the data at once or must you worry about streaming data?
 - Can you bound the maximum size of the input (e.g., to fit in memory)?
- When reading description, identify who is doing what to whom?
 - What are the nouns and verbs used in the description?
 - The nouns are usually data sources or data elements
 - The verbs are often operations you need to perform
 - Look for keywords like min, max, average, median, sort, argmax, sum, find, search, collect, filter out, select, compute, etc...



More advice

- Clearly identify:
 - the source and format of data (same machine? http? xml? csv?)
 - the operation or computation
 - the expected result
 - the output location (disk, remote machine, ...) and format
- Choose simplest possible algorithm & implementation that'll work
 - At first, ignore performance
 - Then worry about getting it into the performance constraints specified

2. Solve



Key ideas for solving problems

- Solving the problem has nothing to do with the computer
- You might not even be asked to code the solution
- If you can't walk through a correct sequence of operations **by hand** on paper, no amount of coding skill will help you!
- All the good programmers I know keep a notepad next to their computers, and it is full of boxes, bubbles, arrows, and notes
- It helps to use established **patterns**, templates, strategies, and common data transformation operations as a crutch

Strategies for solving problems

1. *Start with the end result and work your way backwards*
 - Ask what the prerequisites are for each step
 - The processing step or steps preceding step i compute the data or values needed by step i
 - E.g., median: to pick middle value, previous step must sort data
2. *Reduce or simplify a new problem to a variation of an existing problem with a known solution*
 - Ask what the difference is between the problem you're trying to solve and other problems for which you have a solution
 - E.g., Engineers building a new suspension bridge do not proceed as if such a thing has never been built before

Requisite mathematician joke

- *“A physicist and a mathematician are sitting in a faculty lounge. Suddenly, the coffee machine catches on fire. The physicist grabs a bucket and leaps towards the sink, fills the bucket with water, and puts out the fire.*

Second day, the same two sit in the same lounge. Again, the coffee machine catches on fire. This time, the mathematician stands up, gets a bucket, hands the bucket to the physicist, thus, reducing the problem to one with a known solution.”

Steps in “solve” phase

A. Explore

- Look at the input-output example and imagine how you can manually operate on the input to get the output (using fingers on paper)
- Attempt any manual sequence of operations that appears to be in the right direction, even if you know it's not quite right
- Exploration helps you understand the problem and will trigger more questions, so ask questions
- Write down what you know; e.g., in a geometry problem, you might know which angles or arcs are the same length or that you can make two triangles out of the given rectangle by drawing a line

Steps in “solve” phase continued

B. Reduce

- Can you reduce the problem to known solution by preprocessing the input a bit?

C. Reuse

- Look for and reuse familiar programming patterns like vector sum, min, sort, map, filter, and find
- E.g., to sort a list of numbers (slowly), repeatedly pull then delete the minimum value out of one array and add it to the end of another.

D. Systematize

- Simplify and organize the steps in your process as pseudo-code
- This is your algorithm

Steps in “solve” phase

E. Verify algorithm / process

- Check that your algorithm solves the main problem and the edge cases. Check your algorithm's complexity (performance as function of input size)
- If it's not good enough for the interviewer's constraints, identify the key loops or operations fundamental to your algorithm's complexity
- Iterate on this problem-solving process to reduce complexity
- E.g., can you get rid of a factor of n by converting a linear search to a hash table lookup?

3. Translate your algorithm to code



- A. Write a function that takes your input as parameter(s)
 - Return value of function will typically be the expected problem result
- B. Write a main script that:
 - acquires the data
 - passes it to your function
 - sends the results to the appropriate file or standard output
- C. Translate the algorithm steps into statements in your function
(It's okay if you create helper functions)

4. Verify/test

- Test your code on the representative examples you identified early on in this process
- Now, try some edge cases, which will likely break your code
- Go back to the algorithm and process design phase and alter it to handle the edge cases
- Translate the changes to code
- Verify that you did not break the representative examples and then test on the edge cases

Unit tests

- In a job situation, you'd encode these tests as “unit tests”
- These tests are reproducible and should check edge cases, representative examples, and examples that should fail or cause exceptions
- All code changes over time, which can introduce bugs
- These tests are your primary line of defense against the introduction of bugs in working code (so-called "regressions")
- This is the difference between an amateur and a professional programmer; you cannot safely change code without tests that check the sanity of your system
- For machine learning scripts that just develop models, this might be less true, but it is very true for large or complex systems

What to do when you get stuck

- First thing: Identify exactly what you don't know how to do
- Identifying the tricky bit is a skill that interviewer should look for
- It's a good idea to express verbally, “*Ah. This is what makes this problem tricky*”
- The interviewer might be waiting for you to ask for a hint because they've given you a challenging problem and want to see how you work through it

An example from a real interview

- Computing the median is straightforward for an array sitting on a single machine (e.g., sort and get middle value)
- But, what about data spread across multiple machines?
- Identifying that you can't just take the median of the remote subarray medians is a key part of the interview process
- The solution is tricky and they want to see how fast you can take their hints and come up with a solution

More hints if you're stuck

- What would make this problem easier?
- Try to convert your problem to this easier version by preprocessing the input
- Failing that, solve the simpler version and then work on the harder, more general case
- Multiple failed attempts is part of the game because interviewers won't ask trivial problems, except perhaps during an initial phone screen